

New York City Airbnb Open Data

Francesco Di Salvo
Politecnico di Torino
s282417
s282414@studenti.polito.it

Abstract—In this report we will explain an approach to the New York City Airbnb Open Data, where the aim is to predict the price of a listing, based on fifteen parameters. In particular, the current approach is based on a specific domain analysis, in which it has been discovered the parameters that most affect the price value. So, once the data were cleaned, we were able to accomplish the regression task.

I. PROBLEM OVERVIEW

The current competition is a regression problem based on the New York City Airbnb Open Data, a collection of Airbnb listings in New York City, for the year 2019. The goal of the competition is to predict the price for each insertion. The dataset we have been given was divided into two part:

- *development set* : it contains 39116 records with the price label
- *evaluation set* : it contains 9779 records without the price label

So, we will use the development set for building the regression model, by considering its performance to varying of different hyperparameters. In order to avoid overfitting, we have to explore just the development set, because otherwise the model will not generalize enough. Firstly, it is important to understand the data we are dealing with.

The attributes *id* and *host_id* are incremental values that only identifies the listing and the host, whereas *name* and *host_name* are the related nominal variables. Then, there are 4 features that identify the location of the house. *neighbourhood* and *neighbourhood_group* identify a generic location. The first one is a generalization of second one, in fact there are 221 neighbourhoods and 5 neighbourhood groups: Brooklyn, Queens, Manhattan, Bronx and Staten Island. As it can be clearly seen in Figure 1, there is an overall homogeneous distribution of listing among these neighborhood groups, except for Staten Island that seems to have less apartments for rent. Then, *latitude* and *longitude* represents the exact location of the house.

Another categorical feature is *room_type* that is related to the kind of rent and it can be: *Entire home/apt*, *Private room* and *Shared room*. The first two categories are well distributed in New York City, whereas the last one is less common to see. Figure 2 shows the distribution of these three categories in the city and Figure 3 points out the difference among their frequencies.

There are several other self explanatory attributes, such as *price* (per night), *minimum_nights*, *number_of_reviews*, *last_review* (date) and *reviews_per_month*. These last three

attributes present some null values, in fact if an apartment has no reviews, the attributes last review and the reviews per month will be null too. This issue will be managed during the preprocessing phase. Finally, there is *calculated_host_listings_count* that defines the number of listings for the current host and *availability_365* that represents the availability for the current year. It can be 0 and we can assume that it happens when the host set the house as "unavailable".



Fig. 1. Distribution of listings for each neighbourhood group



Fig. 2. Distribution of listings for each kind of room




Fig. 3. Frequency of type_room

II. PROPOSED APPROACH

A. Data preprocessing

We assumed that *name* *name* *host_name* do not affect the price of a listing, even tho a consistent answer would require a proper text mining analysis. *last_review* is a date object and in this context it could be difficult to manage. So, as a first step, we dropped these three columns. Immediately after, we managed the null values on the reviews fields. In particular, since *null* means that there are no reviews yet, a possible way to manage them is by fill these values with zero.

Then, a critical issues regards the categorical values, in fact any regression model cannot manage nominal values. In particular we are talking about *neighbourhood*, *neighbourhood_group* and *room_type*. An encoding one to one is made by the *factorize()* method provided by pandas, that converts all the distinct strings into distinct values. In order to have a perfect correspondence among the evaluation set and the test set, we used that method on the concatenation of both evaluation and development set. After the encoding, it was possible to distinguish them because the evaluation set had null values on the price label.

Finally, it was possible to look for the linear correlation for the price. As we may imagine, the feature that mostly affects the price is *room_type*. Notice that also a meaning value, for example the *longitude* has a linear correlation of -0.148891. It means that the price has a slight tendency to go down when you go north. Table II defines in detail all the labels and their linear correlation.

Actually I tried to drop the closest feature to 0, that was *id*, but I noticed an overall decline in performance, so I maintained the same structure without dropping any other column.

B. Model selection

The following algorithms have been tested:

- Linear Regression : it is a model that assumes a linear relationship among its features. The prediction is made

TABLE I
ABSOLUTE VALUE OF PRICE CORRELATION

Label	Price correlation
price	1.000000
room_type	0.204980
availability_365	0.082667
calculated_host_listings_count	0.055070
minimum_nights	0.044238
latitude	0.031274
host_id	0.015168
id	0.009273
neighbourhood	-0.013436
number_of_reviews	-0.048254
neighbourhood_group	-0.049987
reviews_per_month	-0.053285
longitude	-0.148891

TABLE II
HYPERPARAMETERS USED FOR THE GRIDSEARCHCV

Model	Hyperparameters	Values
RandomForestRegressor	<i>n_estimators</i> <i>max_features</i>	{300,500,700} {'auto','sqrt','log2'}
LinearRegression	<i>fit_intercept</i> <i>normalize</i>	{True,False} {True,False}

by computing a weighted sum of the input features and the intercept (formally called bias).

- Random Forest Regressor : it is an ensemble of Decision Trees, each one trained with different records and features. It is not so interpretable as a decision tree because of its complexity, but it is much more accurate. In fact, by training several times the model with different features and informations, it is more probable that the noise as well as the overfitting will be reduced.

For both regressors, the best hyperparameters configurations was defined through a grid search, explained in detail in the following section.

C. Hyperparameters tuning

There have been defined two sets of different hyperparameters, defined in Table II.

The grid search trains the model with all the possible configurations and measure its performance by using the cross validation method. This method randomly split the entire dataset into k folds, train the model on k-1 folds and test it with the remaining one. This process is repeated until each fold has been tested. At the end of the process, it is possible to select the best configuration (the one that performs the best score) and use it for testing the evaluation test.

D. Results

The leaderboard is based on the *R2_score*, so we measured our performances with that one. The best configuration for *LinearRegression* was {*fit_intercept*=True, *normalize*=False} and it achieved a result of 0.07815. Instead, the *RandomForestRegressor* obtained an overall performance of 0.20900 with {*max_features*='sqrt', *n_estimators*=700}.

The random forest outperformed the linear regression. Hence, we trained the random forest on the entire development set and we predicted the evaluation set label. The public score obtained is 0.281. In order to compare the hyperparameters tuning and the feature selection phase, we trained the random forest with the default hyperparameters and it obtained a public score of 0.202.

Since the development and the evaluation set initially belonged to the same dataset, we can expect that both price distributions are quite similar over the plane (longitude,latitude). In fact, in Figure 4 there is the price distribution of the development set, and in Figure 5 there is the predicted price distribution from the evaluation set. It can be clearly seen that they are equal.

quite improbable to reach such an accurate model. So the overall result can be considered a good one.

In order to improve the current model, it would be appropriate to do some further analysis on the text fields, such as *name* and *host_name*. It would be interesting to find the most common words in the most reviewed listings and see if there could be a possible relationship among the name of a listing and its popularity, and as a side effect, a rise in the price.




Fig. 4. Price distribution - Development set




Fig. 5. Predicted price distribution - Evaluation set

E. Discussion

The current approach performed much better than the proposal naive solution with the default random forest or with a simple linear regression. Since we are dealing with continuous variables and a huge amount of parameters, it is